

NetGraf: A Collaborative Network Monitoring Stack for Network Experimental Testbeds

Divneet Kaur
dikaur@ucsd.edu
Department of Electrical and
Computer Engineering,
University of California San Diego,
San Diego, CA, USA

Bashir Mohammed (advisor)
bmohammed@lbl.gov
Computing Research Division,
Lawrence Berkeley National
Laboratory,
Berkeley, CA, USA

Mariam Kiran (advisor)
mkiran@es.net
Scientific Networking Division,
Lawrence Berkeley National
Laboratory,
Berkeley, CA, USA

ABSTRACT

Network performance monitoring collects heterogeneous data such as network flow data to give an overview of network performance, and other metrics, necessary for diagnosing and optimizing service quality. However, due to disparate and heterogeneity, to obtain metrics and visualize entire data from several devices, engineers have to log into multiple dashboards.

In this paper we present NetGraf, a complete end-to-end network monitoring stack, that uses open-source network monitoring tools and collects, aggregates, and visualizes network measurements on a single easy-to-use real-time Grafana dashboard. We develop a novel NetGraf architecture and can deploy it on any network testbed such as Chameleon Cloud by single easy-to-use script for a full view of network performance in one dashboard.

This paper contributes to the theme of automating open-source network monitoring tools software setups and their usability for researchers looking to deploy an end-to-end monitoring stack on their own testbeds.

KEYWORDS

network monitoring tools, real-time dashboards, deployable solution

ACM Reference Format:

Divneet Kaur, Bashir Mohammed (advisor), and Mariam Kiran (advisor). 2020. NetGraf: A Collaborative Network Monitoring Stack for Network Experimental Testbeds. In *SC '20: ACM/IEEE Supercomputing Conference, November 17–19, 2020, Online*. ACM, 3 pages. <https://doi.org/10.1145/nnnnnnnn>

1 INTRODUCTION

Cloud infrastructure is composed of heterogeneous resources made up of hardware, virtualization, storage, and networking components [1]. Network performance monitoring (NPM) is the process of visualizing, monitoring, optimizing, troubleshooting and reporting on the service quality of your network as experienced by your users [2]. Different NPM tools collect different data such as packet

loss, network flow which when combined provides a complete picture of the network infrastructure. This helps monitor and analyze a network's performance, availability, and other important metrics. But this heterogeneity comes with a challenge when network engineers try to visualize all the network metrics coming from different monitoring tools without single dashboard.

We develop NetGraf, a collaborative network monitoring stack that provides a holistic view of the network system providing visualizations of various metrics from different monitoring tools in a single dashboard for valuable insight on the network. While developing monitoring stack, our main contributions are:

- We explore six diverse open-source network monitoring tools and package them for any network infrastructure by developing a monitoring pipeline that fetches data from these tools, aggregates them and stores in a time-series database [3].
- We develop an Application Programming Interface (API) to define interactions between these tools and Grafana, an open-source visualization software, in order to generate visualizations of collected metrics and network statistics.

2 METHODOLOGY: NETGRAF ARCHITECTURE

Figure 1 presents the NetGraf architecture. NetGraf is designed for supporting multiple network monitoring tools, for different network devices such as large scale hardware infrastructure and application servers. It supports seamless plug-in for new NPM tools. It also allows monitoring, collecting, identifying and visualizing a wide spectrum of network data in a single dashboard that enables network engineers to identify mishaps like degradation and outages that occur in a network experimental testbed.

The architecture consists of three main modules:

2.0.1 Network and Application Module. The network topology deployed on Chameleon testbed with monitoring tools.

2.0.2 Data Collector and Aggregator Module. Various monitoring tools were installed to collect metrics and network statistics, such as ntopng and netdata (Table 1). Since Prometheus can scrape metrics from multiple nodes, it is installed on one node and scrapes metrics from all other nodes. Zabbix was installed on one node collecting server related metrics. perfSONAR, a network measurement toolkit, was installed on Chicago and Texas sites.

These tools were connected to databases for storage - ntopng, netdata were connected to InfluxDB, a time-series database. Prometheus,

Permission to make digital or hard copies of all or part of this work for personal or commercial use, by users registered with ACM, is granted by ACM Publishing Department for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SC '20, November 17–19, 2020, Online
© 2020 Association for Computing Machinery.
ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00
<https://doi.org/10.1145/nnnnnnnn>

2020-09-17 01:49. Page 1 of 1–3.

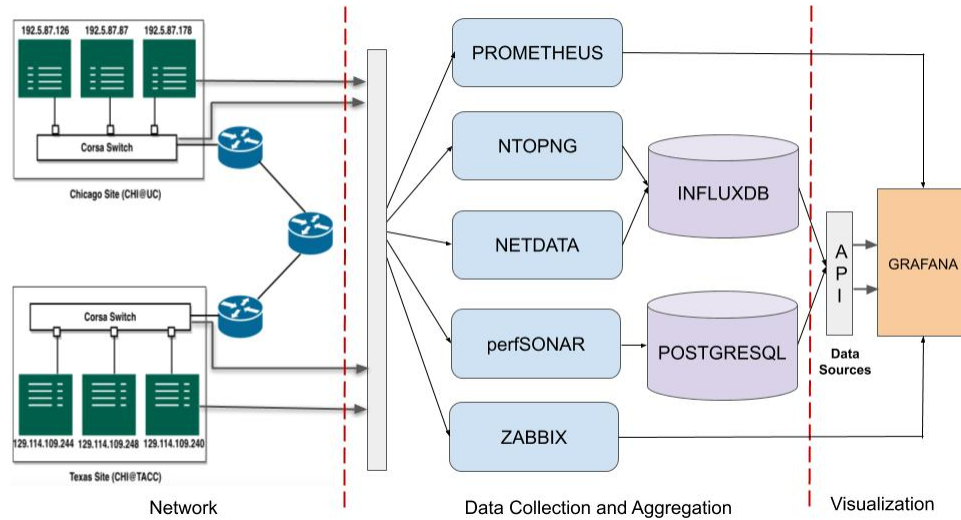


Figure 1: NetGraf architecture composes of three main modules: Network, Data collection and Aggregation, Visualization module.

Table 1: Chameleon Core Network Environment settings.

Node Specification	Network Monitoring Tools on Chameleon testbed					
	Prometheus	ntopng	netdata	perfSONAR	Zabbix	Grafana
IP Address	192.168.100.11	Installed on all nodes	Installed on all nodes	192.168.100.16	192.168.100.13	192.168.100.14
Floating IP	192.5.87.178	Installed on all nodes	Installed on all nodes	192.5.87.157	192.5.87.126	192.5.87.126
Gateway	192.168.100.1	192.168.100.1	192.168.100.1	192.168.100.1	192.168.100.1	192.168.100.1
Listening Port	9090	3000	19999	861	10050	3000
Switches	Corsa Switch	Corsa Switch	Corsa Switch	Corsa Switch	Corsa Switch	Corsa Switch
OS	CC-Ubuntu18.04	CC-Ubuntu18.04/16.04, Centos7	CC-Ubuntu18.04/16.04, Centos7	CC-Ubuntu18.04	CC-CentOS7	CC-Ubuntu18.04

Zabbix have an inbuilt database. perfSONAR's collected results were archived in a relational database, postgresSQL.

2.0.3 Monitoring and Visualization Module. To generate visualizations from the available metrics, we created an API between the databases and Grafana. This was established by adding Influxdb, Postgresql, Prometheus and Zabbix as *data sources* in Grafana.

Since a large number of metrics were being collected, we performed an elimination process where network metrics such as Transmission Control Protocol, throughput and loss were selected. This helped us create a dashboard with relevant metrics only.

3 RESULTS: LESSONS LEARNT

Figure 2 shows a snapshot of five nodes from Chameleon testbed located at Chicago for 3 hours.

While identifying efficient routes to connect the monitoring tools to Grafana, we explored two other approaches,

- The monitoring tools were connected to Prometheus which recorded only node metrics such as disk storage, which was not network data.

- The data was directly fed to Grafana using plugins. Due to lack of direct plugins for all tools and databases for storage, this approach was not efficient as well.

4 CONCLUSION AND FUTURE WORK

Monitoring and understanding network infrastructure performance is essential to learn network experimentation. In this work, we present a unique monitoring approach which can collect and store network metrics and solve the heterogeneity of diverse network monitoring tools by displaying them all in a single dashboard. In future, we will apply machine learning algorithms to the dashboard to provide more insights on network performance and availability analysis.

ACKNOWLEDGMENTS

We would like to thank Paul Ruth for the technical support. This work was supported by U.S. Department of Energy, Office of Science Early Career Research Program for 'Large-scale Deep Learning for Intelligent Networks' Contract no FP00006145.

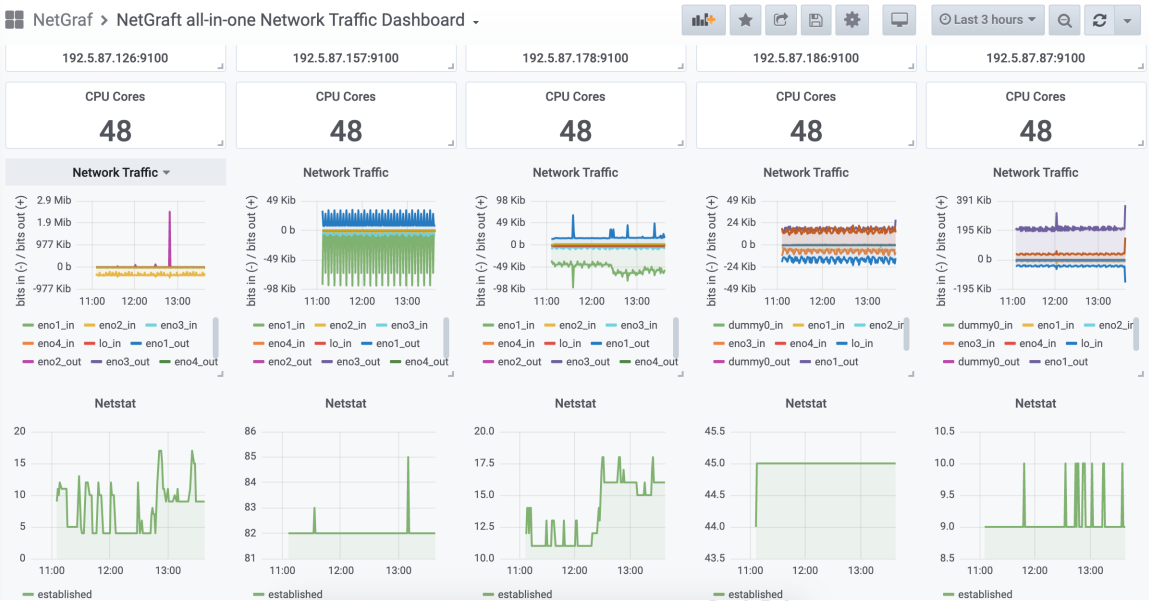


Figure 2: NetGraf showing a snapshot of network metrics on Grafana.

REFERENCES

[1] Dan C Marinescu. 2017. *Cloud computing: theory and practice*. Morgan Kaufmann.

[2] Srinivas Narayana, Anirudh Sivaraman, Vikram Nathan, Prateesh Goyal, Venkat Arun, Mohammad Alizadeh, Vimalkumar Jeyakumar, and Changhoon Kim. 2017. Language-directed hardware design for network performance monitoring. In *Proceedings of the Conference of the ACM Special Interest Group on Data Communication*. 85–98.

[3] Salman Taherizadeh, Andrew C Jones, Ian Taylor, Zhiming Zhao, Paul Martin, Vlado Stankovski, et al. 2016. Runtime network-level monitoring framework in the adaptation of distributed time-critical cloud applications. In *The 22nd International Conference on Parallel and Distributed Processing Techniques and Applications*.