


# 1 Multi-agent Markov Decision Process

The package routing problem can be formulated as a multi-agent Markov decision process. Let  $n$  denote the number of routers in the network. A multi-agent Markov decision process is characterized by a tuple  $M = \langle S, A^i, P, R \rangle$  where  $S$  is the global state space shared by all the routers,  $A^i$  is the set of actions that router  $i$  can execute. Let  $A = \prod_{i=1}^n A^i$  be the joint action space of all the routers.  $P : S \times A \times S \rightarrow [0, 1]$  is a state transition probability and  $R : S \times A \rightarrow R$  is the shared reward function. Each router has a policy  $\pi_i$  representing the probability of choosing action  $a^i$  at state  $s_t$ . The policy  $\pi_i$  is approximated by a parameterized function  $\pi_{\theta_i}^i$ . Let  $\pi_\theta = \prod_{i=1}^n \pi_{\theta_i}^i$  be the joint policy of all the routers. The objective of the routers is to find the optimal policy parameters that maximize the expected reward, which is given by  $J(\theta)$

$$J(\pi_\theta) = E_{x \sim d_{\pi_\theta}, a_i \sim \pi_{\theta_i}^i} [\pi(s, a_1, \dots, a_n) R(x, a)], \quad (1)$$

where  $d_{\pi_\theta}$  is the stationary distribution of the Markov chain under policy  $\pi_\theta$ .

**In summary, both state and reward are global and shared by all the routers. Action is private. Is it possible to only use local state or reward?**



## 2 Algorithm

### 2.1 MAML

Model-agnostic meta reinforcement learning is one kind of meta RL algorithm which achieves good generalization performance on new tasks by updating policy model parameters. As shown in Figure 1, a good policy model parameter  $\theta$  should be close to all the optimal parameters of the tasks which makes  $\theta$  the best parameters initialization that can quickly adapt to different new tasks. When we have, for example, three different new tasks 1, 2 and 3, a gradient step is taken for each task (the gray lines). We can see that the parameters are close to all the 3 optimal parameters of task 1, 2, and 3 which makes  $\theta$  the best parameters initialization that can quickly adapt to different new tasks.

As a result, only a small change in the parameters  $\theta$  will lead to an optimal minimization of the loss function of any task.

## 2.2 Trust Region Policy Optimization

TRPO algorithm limits the parameter changes (policy changes).

$$\begin{aligned} \max \quad & E_{s \sim d_{\theta_{\text{old}}}, a_i \sim \pi_{\theta_{\text{old}}}^i} \left[ \frac{\pi_{\theta}(a_1, \dots, a_n | s)}{\pi_{\theta_{\text{old}}}(a_1, \dots, a_n | s)} A_{\pi_{\theta_{\text{old}}}}(s, a_1, \dots, a_n) \right] \\ \text{s.t.} \quad & E_{s \sim d_{\theta_{\text{old}}}} [D_{KL}(\pi_{\theta_{\text{old}}}^i(\cdot | s) || \pi_{\theta}^i(\cdot | s))] \leq \delta, \quad \forall i \in [1, n], \end{aligned} \quad (2)$$

where  $\pi_{\theta}(a_1, \dots, a_n | s) = \prod_{i=1}^n \pi_{\theta}^i(a_i | s)$ . The solution of Equation (2) can be found in Appendix C in "Trust Region Policy Optimization".

## 2.3 Multi-agent MAML

---

**Algorithm 1:** Multi-agent MAML algorithm

---

**Require:**  $p(\tau)$ : distributions of tasks;  
**Require:**  $\alpha$ : step size hyper-parameter;  
randomly initialize  $\theta_i, i \in [1, n]$ ;  
**while** *not done* **do**  
    sample batch of tasks  $\tau_i \sim p(\tau)$ ;  
    **for all**  $\tau_i$  **do**  
        Sample  $K$  trajectories  $D = \{(x_1, a_1, r_1, x_2, \dots, x_H)\}$  using  $\pi_{\theta_i}, i \in [1, n]$  in  $\tau_i$ ;  
        Evaluate  $\nabla_{\theta_i} J(\pi_{\theta})$  using  $D$  and  $J(\pi_{\theta})$  in Equation (1);  
        Compute adapted parameters with gradient descent:  
         $\theta'_i = \theta_i - \alpha \nabla_{\theta_i} J(\pi_{\theta})$ ;  
        Sample  $K$  trajectories  $D'_i = \{(x_1, a_1, r_1, x_2, \dots, x_H)\}$  using  $\pi_{\theta'_i}, i \in [1, n]$  in  $\tau_i$ ;  
    **end**  
    Update  $\theta$  according to  $D'_i$  and Equation (2)  
**end**

---